

## # Lab: Logistic Regression LDA QDA KNN

### #4.6.1 The Stock Market data

```
In [89]: 1 library(ISLR)
          2 names(Smarket)

'Year' 'Lag1' 'Lag2' 'Lag3' 'Lag4' 'Lag5' 'Volume' 'Today' 'Direction'
```

```
In [90]: 1 dim(Smarket)

1250 9
```

In [91]:

```
1 summary(Smarket)

      Year      Lag1      Lag2      Lag3
Min.   :2001  Min.   :-4.922000  Min.   :-4.922000  Min.   :-4.92
2000
1st Qu.:2002  1st Qu.: -0.639500  1st Qu.: -0.639500  1st Qu.: -0.64
0000
Median :2003  Median :  0.039000  Median :  0.039000  Median :  0.03
8500
Mean   :2003  Mean    :  0.003834  Mean    :  0.003919  Mean    :  0.00
1716
3rd Qu.:2004  3rd Qu.:  0.596750  3rd Qu.:  0.596750  3rd Qu.:  0.59
6750
Max.   :2005  Max.    :  5.733000  Max.    :  5.733000  Max.    :  5.73
3000

      Lag4      Lag5      Volume      Today
Min.   :-4.922000  Min.   :-4.92200  Min.   :0.3561  Min.   :-4.9
22000
1st Qu.: -0.640000  1st Qu.: -0.64000  1st Qu.:1.2574  1st Qu.: -0.6
39500
Median :  0.038500  Median :  0.03850  Median :1.4229  Median :  0.0
38500
Mean   :  0.001636  Mean    :  0.00561  Mean    :1.4783  Mean    :  0.0
03138
3rd Qu.:  0.596750  3rd Qu.:  0.59700  3rd Qu.:1.6417  3rd Qu.:  0.5
96750
Max.   :  5.733000  Max.    :  5.73300  Max.    :3.1525  Max.    :  5.7
33000
Direction
Down:602
Up   :648
```

In [92]:

```
1 #cor(Smarket)
2 #generating error, because Dir is not numremic
3
4 #lag1-lag5- %return of previous trading days
5 #Volume-No. of shares traded of previous days in billion
6 #Today- %age return on the date in question
7 #Direction- whether market was up or down
8 ?Smarket
```

In [93]:

```

1 cor(Smarket[, -9])
2 #produces a matrix that contains all of the pairwise
3 #correlations among the predictors in a data set

```

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	
<b>Year</b>	1.00000000	0.029699649	0.030596422	0.033194581	0.035688718	0.029787995	0.5
<b>Lag1</b>	0.02969965	1.00000000	-0.026294328	-0.010803402	-0.002985911	-0.005674606	0.0
<b>Lag2</b>	0.03059642	-0.026294328	1.00000000	-0.025896670	-0.010853533	-0.003557949	-0.0
<b>Lag3</b>	0.03319458	-0.010803402	-0.025896670	1.00000000	-0.024051036	-0.018808338	-0.0
<b>Lag4</b>	0.03568872	-0.002985911	-0.010853533	-0.024051036	1.00000000	-0.027083641	-0.0
<b>Lag5</b>	0.02978799	-0.005674606	-0.003557949	-0.018808338	-0.027083641	1.00000000	-0.0
<b>Volume</b>	0.53900647	0.040909908	-0.043383215	-0.041823686	-0.048414246	-0.022002315	1.0
<b>Today</b>	0.03009523	-0.026155045	-0.010250033	-0.002447647	-0.006899527	-0.034860083	0.0

In [94]:

```


1 #observation:correlations between the lag variables and
2 #today's returns are close to zero
3 #little correlation between today's returns and previous days'(vol

```

```
In [95]: 1 attach(Smarket)
2         plot(Volume)
3
4         #vol is increasing
```

The following objects are masked from Smarket (pos = 6):

Direction, Lag1, Lag2, Lag3, Lag4, Lag5, Today, Volume, Year



## # 4.6.2 Logistic regression

```
In [96]: 1 glm.fits=glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume,data=Smarke
2
3 # glm() function fits generalized linear models,
4 #a class of models that includes logistic regression
```

```
In [97]: 1 summary(glm.fits)
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
     Volume, family = binomial, data = Smarket)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.446	-1.203	1.065	1.145	1.326

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-0.126000	0.240736	-0.523	0.601
Lag1	-0.073074	0.050167	-1.457	0.145
Lag2	-0.042301	0.050086	-0.845	0.398
Lag3	0.011085	0.049939	0.222	0.824
Lag4	0.009359	0.049974	0.187	0.851
Lag5	0.010313	0.049511	0.208	0.835
Volume	0.135441	0.158360	0.855	0.392

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1731.2 on 1249 degrees of freedom

Residual deviance: 1727.6 on 1243 degrees of freedom

AIC: 1741.6

Number of Fisher Scoring iterations: 3

```
In [98]: 1 #observations
2 #lag1 has min p-value,
3 #negative coefficient for this predictor suggests that-
4 #if the market had a positive return yesterday,
5 #then it is less likely to go up today
6
7 #however, 0.15 is still large value of p
8 # so there is no clear evidence of a real association
9 #between Lag1 and Direction.
```

```
In [99]: 1 #coef() function in order to access the
          2 #coefficients for this fitted model.
          3 coef(glm.fits)
```

```
(Intercept) -0.126000256559266
Lag1 -0.0730737458900261
Lag2 -0.0423013440073083
Lag3 0.0110851083796762
Lag4 0.0093589383702787
Lag5 0.0103130684758179
Volume 0.13544065885916
```

```
In [100]: 1 summary(glm.fits)$coef
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-0.126000257	0.24073574	-0.5233966	0.6006983
Lag1	-0.073073746	0.05016739	-1.4565986	0.1452272
Lag2	-0.042301344	0.05008605	-0.8445733	0.3983491
Lag3	0.011085108	0.04993854	0.2219750	0.8243333
Lag4	0.009358938	0.04997413	0.1872757	0.8514445
Lag5	0.010313068	0.04951146	0.2082966	0.8349974
Volume	0.135440659	0.15835970	0.8552723	0.3924004

```
In [101]: 1 summary(glm.fits)$coef[,4]
```

```
(Intercept) 0.600698319413355
Lag1 0.145227211568647
Lag2 0.398349095427021
Lag3 0.824333346101536
Lag4 0.851444506926455
Lag5 0.834997390499829
Volume 0.392400433202429
```

```
In [102]: 1 #The predict() function can be used to predict the probability tha
2 #the market will go up, given values of the predictors
3 #The predict() function is used similiary to generate predictions
4 #for the response variable.
5 glm.probs=predict(glm.fits,type="response")
6 glm.probs[1:10] #printed only the first 10 probabilities
```

```
1 0.507084133395402
2 0.481467878454591
3 0.481138835214201
4 0.515222355813022
5 0.510781162691538
6 0.506956460534911
7 0.492650874187038
8 0.509229158207377
9 0.517613526170958
10 0.488837779771376
```

```
In [103]: 1 contrasts (Direction)
2 #Use the contrasts() function to see the dummy variables generated
3 #for values in the categorical variable Direction.
4
5 #In order to make a prediction as to whether the market will go
6 #up or down on a particular day, we must convert these predicted
7 #probabilities into class labels, Up or Down.
```

	Up
Down	0
Up	1

```
In [104]: 1 glm.pred=rep("Down",1250) #creates a vector of 1,250 Down element
2 glm.pred[glm.probs >.5]="Up" #transforms to Up all of the elements
3
4 #creating a vector of class predictions based on whether
5 #the predicted probability of a market increase is
6 #greater than or less than 0.5
7
```

```
In [105]: 1 #We can generate a confusion matrix between the predicted direction
          2 #and the actual direction from the variable Direction
          3 #using the table()function.
          4
          5 table(glm.pred,Direction) #confusion matrix
```

```
      Direction
glm.pred Down Up
Down   145 141
Up     457 507
```

```
In [106]: 1 #diagonal elements of the confusion matrix indicate correct predictions
          2 #while the off-diagonals represent incorrect predictions.
          3 #market would go up on 507 days and that it would go down on 145 days
          4
          5 (507+145) /1250 #correct prediction
          6
          7 #logistic regression correctly predicted the movement of the market
          8 #training error--> 100 - 52.2 = 47.8 %
```

```
0.5216
```

```
In [107]: 1 #We then divide our dataset into training set and test set.
          2 #The training set will include observations from 2001-2004
          3 #and the test set from the year 2005.
          4
          5 train=(Year <2005) #train is a vector of 1,250 elements as observations
          6 Smarket.2005= Smarket[! train ,] #pick submatrix of market data, but
          7 dim(Smarket.2005)
```

```
252 9
```

```
In [108]: 1 Direction.2005=Direction[!train]
          2
```

```
In [109]: 1 glm.fits=glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume , data=Smarket.2005)
```

```
In [110]: 1 glm.probs=predict(glm.fits,Smarket.2005,type="response")
```

```
In [111]: 1 #we have trained and tested our model on two completely separate-
          2 # training was performed using only the dates before 2005,
          3 #and testing was performed using only the dates in 2005.
```



```
In [112]: 1 glm.pred=rep("Down",252)
2          glm.pred[glm.probs >.5]="Up"
3          table(glm.pred,Direction.2005)
4
5          #To improve the predictive performance, we can restrict the predict
6          #to only those with the strongest relationship to the response var
7          #In this case, we limit the variables to Lag1 and Lag2.
8
```

```
          Direction.2005
glm.pred Down Up
Down      77 97
Up        34 44
```

```
In [113]: 1 mean(glm.pred==Direction.2005)

0.48015873015873
```

```
In [114]: 1 mean(glm.pred!=Direction.2005)

0.51984126984127
```

```
In [115]: 1 glm.fits=glm(Direction~Lag1+Lag2,data=Smarket ,family=binomial, su
```

```
In [116]: 1 predict(glm.fits,newdata=data.frame(Lag1=c(1.2,1.5),
2          Lag2=c(1.1,-0.8)),type="response")



1 0.479146239171912
2 0.496093872956532
```

## # Linear Discriminant analysis

```
In [117]: 1 library(MASS)
```

```
In [118]: 1 lda.fit=lda(Direction~Lag1+Lag2,data=Smarket ,subset=train)
```

In [119]: 1 plot(lda.fit)



```
In [120]: 1 lda.pred = predict(lda.fit, Smarket.2005)
          2 names(lda.pred)
          3
          4 #The predict() function for an LDA model returns
          5 #list of three elements representing the predicted class,
          6 #the posterior probabilities and the linear discriminants
          7 #
```

'class' 'posterior' 'x'

```
In [121]: 1 lda.class = lda.pred$class
          2 table(lda.class, Direction.2005)
          3
          4 #comparing the predicted class with the predicted directed obtained
          5 #from logistic regression reviously and stored in the vector Direc
```

```
          Direction.2005
lda.class Down  Up
      Down   35  35
      Up    76 106
```

```
In [122]: 1 mean(lda.class == Direction.2005)
```

```
0.55952380952381
```

```
In [123]: 1 sum(lda.pred$posterior[, 1] >= 0.5)
```

```
70
```

```
In [124]: 1 sum(lda.pred$posterior[, 1] < 0.5)
```

```
182
```

```
In [125]: 1 lda.pred$posterior[1:20, 1]
          2 #We can inspect the posterior probabilities of the LDA model
          3 #from the posterior vector of the fitted model.
```

```
          999 0.490179249818258
          1000 0.479218499099683
          1001 0.466818479852065
          1002 0.474001069455248
          1003 0.492787663967445
          1004 0.493856154997504
          1005 0.495101564646223
          1006 0.487286099421815
          1007 0.490701348960405
          1008 0.484402624071869
          1009 0.490696276120968
          1010 0.511998846261919
          1011 0.489515226936648
          1012 0.470676122211879
          1013 0.474459285611829
          1014 0.479958339148108
          1015 0.493577529465861
          1016 0.503089377118306
          1017 0.497880612141404
          1018 0.488633086516518
```

```
In [126]: 1 lda.class[1:20]
```

```
Up Up Up Up Up Up Up Up Up Up Up Up Down Up Up Up Up Up
Down Up Up
```

► Levels:

```
In [127]: 1 #We can also set the posterior probabilities to different
          2 #thresholds for making predictions.
```

## # 4.6.4 Quadratic Discriminant Analysis

```
In [128]: 1 qda.fit = qda(Direction ~ Lag1 + Lag2, data = Smarket, subset = tr
2 qda.fit
```

Call:

```
qda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
```

Prior probabilities of groups:

```
      Down      Up
0.491984 0.508016
```

Group means:

```
      Lag1      Lag2
Down 0.04279022 0.03389409
Up   -0.03954635 -0.03132544
```

```
In [129]: 1 #We can make predictions using predict() just as we did for an LDA
2 #and compare them to the results from the logistic regression
3 qda.class = predict(qda.fit, Smarket.2005)$class
4 table(qda.class, Direction.2005)
```

```
      Direction.2005
qda.class Down Up
      Down   30  20
      Up    81 121
```

```
In [130]: 1 mean(qda.class == Direction.2005)
```

```
0.599206349206349
```

## #4.6.5 K-Nearest Neighbors

```
In [131]: 1 #The class package offers a number of classification algorithms
2 #including K-Nearest Neighbors. Before we can run the KNN algorithm
3 #we need to split our dataset into training and test subsets.
4 #After splitting the dataset, the cbind() is used to bind the Lag1
5 #into a matrix for each subset.
6
7 library(class)
8 train.X=cbind(Lag1 ,Lag2)[train ,] # training data,
9 test.X=cbind(Lag1,Lag2)[!train,]  # predictors associated with th
10 train.Direction=Direction [train] #class labels for the training
```

```
In [132]: 1 set.seed(1)
          2 knn.pred = knn(train.X, test.X, train.Direction, k = 1)
          3 table(knn.pred, Direction.2005)
          4
          5 # set.seed() to ensure that repeated runs produce consistent results
          6 #make predictions about the market direction in 2005.
```

```
          Direction.2005
knn.pred Down Up
Down     43 58
Up       68 83
```

```
In [133]: 1 (83 + 43)/252
```

```
0.5
```

```
In [134]: 1 knn.pred <- knn(train.X, test.X, train.Direction, k = 3)
          2 table(knn.pred, Direction.2005)
          3
          4 #can repeat the fit with K = 3.
```

```
          Direction.2005
knn.pred Down Up
Down     48 54
Up       63 87
```

```
In [135]: 1 mean(knn.pred == Direction.2005)
```

```
0.535714285714286
```

## # 4.6.6 An Application to Caravan Insurance Data

In [136]: `1 attach(Caravan) #function to make the Caravan dataset available`

The following objects are masked from Caravan (pos = 4):

```
AAANHANG, ABESAUT, ABRAND, ABROM, ABYSTAND, AFIETS, AGEZONG,
AINBOED, ALEVEN, AMOTSCO, APERSAUT, APERSONG, APLEZIER, ATRACTOR,
AVRAAUT, AWABEDR, AWALAND, AWAOREG, AWAPART, AWERKT, AZEILPL,
MAANTHUI, MAUT0, MAUT1, MAUT2, MBERARBG, MBERARBO, MBERBOER,
MBERHOOG, MBERMIDD, MBERZELF, MFALLEEN, MFGEKIND, MFEWKIND,
MGEMLEEF, MGEMOMV, MGODGE, MGODOV, MGODPR, MGODRK, MHHUUR, MHKOOP
,
MINK123M, MINK3045, MINK4575, MINK7512, MINKGEM, MINKM30, MKOOPKL
A,
MOPLHOOG, MOPLLAAG, MOPLMIDD, MOSHOOFD, MOSTYPE, MRELGE, MRELOV,
MRELSA, MSKA, MSKB1, MSKB2, MSKC, MSKD, MZFONDS, MZPART, PAANHANG
,
PBESAUT, PBRAND, PBROM, PBYSTAND, PFIETS, PGEZONG, PINBOED, PLEVE
N,
PMOTSCO, PPERSAUT, PPERSONG, PPLEZIER, PTRACTOR, PVRAAUT, PWABEDR
,
PWALAND, PWAOREG, PWAPART, PWERKT, PZEILPL, Purchase
```

In [137]: `1 dim(Caravan)`

5822 86

In [138]: `1 summary(Purchase)`

```
      No    5474
      Yes    348
```

In [139]: `1 348/5822`

0.0597732737890759

In [140]: `1 standardized.X <- scale(Caravan[, -86])
2 var(Caravan[, 1])
3
4 #scale() function to scale the dataset with a mean of zero and sta`

165.037847395189

```
In [141]: 1 var(Caravan[, 2])
```

```
0.164707781931954
```

```
In [142]: 1 var(standardized.X[, 1])
```

```
1
```

```
In [143]: 1 var(standardized.X[, 2])
```

```
1
```

```
In [144]: 1 test = 1:1000
2 train.X = standardized.X[-test, ]
3 test.X = standardized.X[test, ]
4 train.Y = Purchase[-test]
5 test.Y = Purchase[test]
6 set.seed(1)
7 knn.pred = knn(train.X, test.X, train.Y, k = 1)
8 mean(test.Y != knn.pred)
```

```
0.118
```

```
In [145]: 1 mean(test.Y != "No")
```

```
0.059
```

```
In [146]: 1 table(knn.pred, test.Y)
```

```
      test.Y
knn.pred No Yes
No      873  50
Yes     68   9
```

```
In [147]: 1 9/(68 + 9)
```

```
0.116883116883117
```

```
In [148]: 1 knn.pred = knn(train.X, test.X, train.Y, k = 3) #takeing k = 3
2 table(knn.pred, test.Y)
```

```
      test.Y
knn.pred No Yes
No      920  54
Yes     21   5
```



In [149]: 1 5/26

0.192307692307692

In [150]: 1 knn.pred = knn(train.X, test.X, train.Y, k = 5) #k = 5  
2 table(knn.pred, test.Y)

```
      test.Y
knn.pred No Yes
      No  930  55
      Yes   11   4
```

In [151]: 1 glm.fit = glm(Purchase ~ ., data = Caravan, family = binomial, sub  
2  
3 *#compare the KNN model with a logistic regression using glm() and*

Warning message:  
"glm.fit: fitted probabilities numerically 0 or 1 occurred"

In [152]: 1 glm.probs = predict(glm.fit, Caravan[test, ], type = "response")  
2 glm.pred = rep("No", 1000)  
3 glm.pred[glm.probs > 0.5] <- "Yes"  
4 table(glm.pred, test.Y)

```
      test.Y
glm.pred No Yes
      No  934  59
      Yes   7   0
```

In [155]: 1 glm.pred = rep("No", 1000)  
2 glm.pred[glm.probs > 0.25] = " Yes"  
3 table(glm.pred, test.Y)

```
      test.Y
glm.pred No Yes
      Yes  22  11
      No  919  48
```

In [156]: 1 11/(22 + 11)

0.333333333333333

