


#4


```
In [1]: 1 set.seed(113)
2 x = rnorm(100)
3 y = 3 * x^2 + 4 + rnorm(100)
4 train = sample(100, 50)
5 y[train] = y[train] + 3
6 y[-train] = y[-train] - 3
7
8 plot(x[train], y[train], pch="+", lwd=4, col="red",
9      ylim=c(-4, 20), xlab="X", ylab="Y")
10 points(x[-train], y[-train], pch="o", lwd=4, col="blue")
```



In [2]:

```
1  #plot is non-linearly seperable,  
2  #creating train and test dfs by taking half of +ve and -ve classes  
3  #to create a new vector  
4  
5  library(e1071)  
6  set.seed(113)  
7  z = rep(0, 100)  
8  z[train] = 1  
9  # Take 25 observations each from train and -train  
10 final.train = c(sample(train, 25), sample(setdiff(1:100, train), 25))  
11 data.train = data.frame(x=x[final.train], y=y[final.train],  
12                        z=as.factor(z[final.train]))  
13 data.test = data.frame(x=x[-final.train], y=y[-final.train],  
14                       z=as.factor(z[-final.train]))
```

```
In [3]: 1 svm.linear = svm(z~., data=data.train, kernel="linear", cost=10)
        2 plot(svm.linear, data.train)
```




```
In [4]: 1 #plot shows linear boundry
```

```
In [5]: 1 table(z[final.train], predict(svm.linear, data.train))
```

```
      0  1
0 17  8
1  2 23
```

```
In [6]: 1 set.seed(32325)
        2 svm.poly = svm(z~., data=data.train, kernel="polynomial", cost=10)
        3 plot(svm.poly, data.train)
```




```
In [7]: 1 table(z[final.train], predict(svm.poly, data.train))
```

```
      0  1
0 13 12
1   2 23
```

```
In [8]: 1 #This is a default polynomial kernel with degree 3.
```

```
In [10]: 1 set.seed(353)
          2 svm.radial = svm(z~., data=data.train, kernel="radial", gamma=1,
          3                   cost=10)
          4 plot(svm.radial, data.train)
```




```
In [11]: 1 table(z[final.train], predict(svm.radial, data.train))
```


```
      0  1
0 25  0
1  0 25
```

```
In [12]: 1 #this classifier perfectly classifies train data.
```

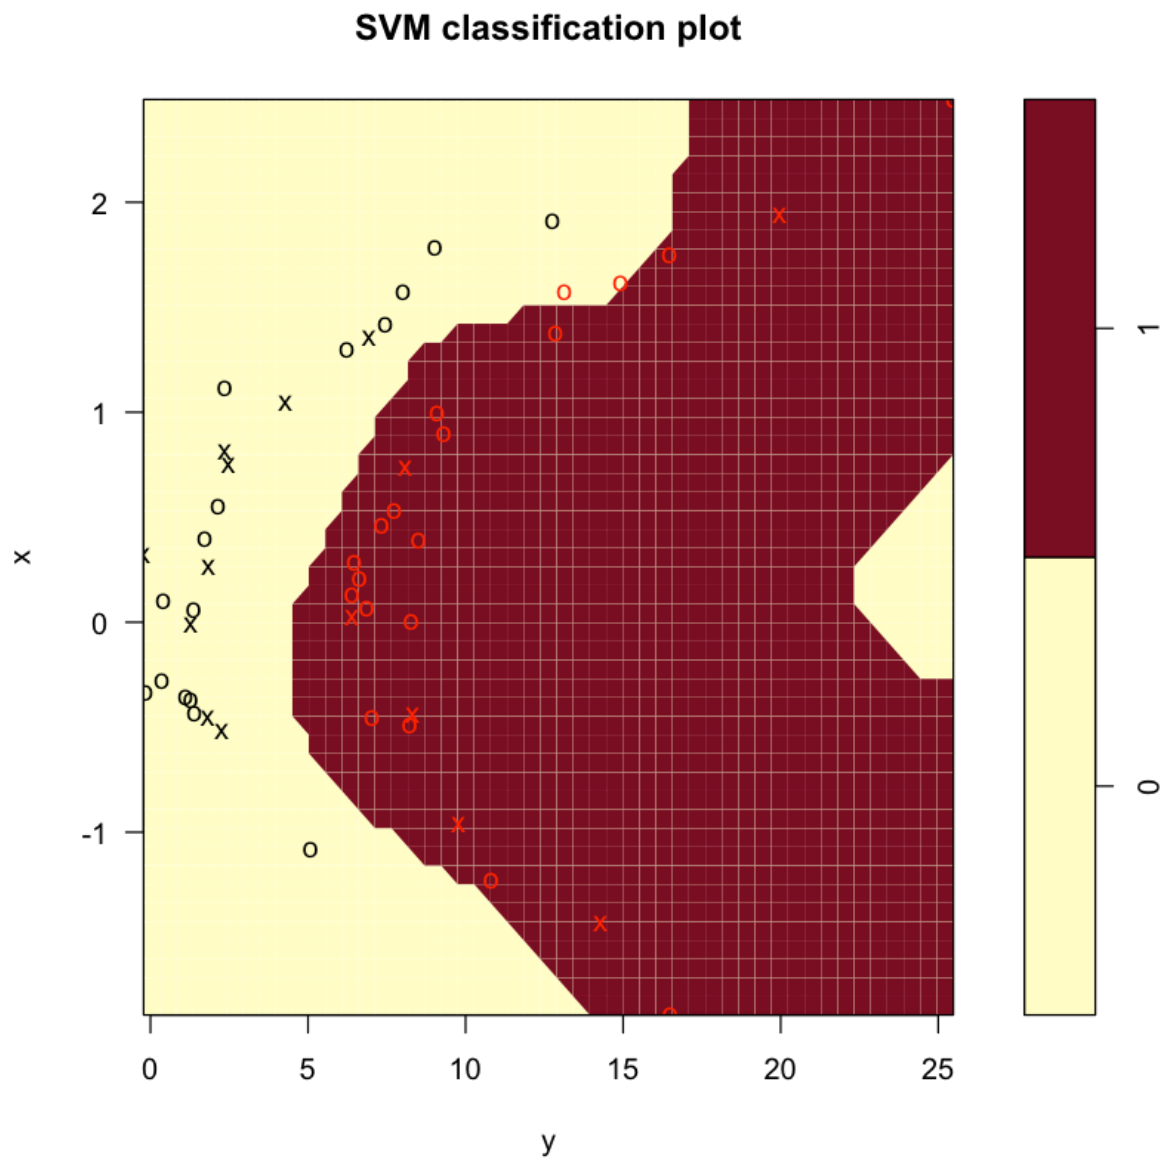
```
In [13]: 1 plot(svm.linear, data.test)
```



```
In [14]: 1 plot(svm.poly, data.test)
```




```
In [15]: 1 plot(svm.radial, data.test)
```



```
In [16]: 1 table(z[-final.train], predict(svm.linear, data.test))
```

0	1
0	18
1	0
25	

```
In [17]: 1 table(z[-final.train], predict(svm.poly, data.test))
```

```
      0  1  
0 11 14  
1   1 24
```

```
In [18]: 1 table(z[-final.train], predict(svm.radial, data.test))
```

```
      0  1  
0 25  0  
1   1 24
```

```
In [19]: 1 #linear, polynomial and radial basis kernels classify  
2 #7, 15, and 1 test points incorrectly respectively.  
3 #Radial kernel is the best and has less test misclassification error  
4 #than linear and polynomial.
```


#5

#a

```
In [20]: 1 set.seed(333)  
2 x1 = runif(500) - 0.5  
3 x2 = runif(500) - 0.5  
4 y = 1 * (x1^2 - x2^2 > 0)
```

#b

```
In [21]: 1 plot(x1[y == 0], x2[y == 0], col = "red", xlab = "X1",  
2           ylab = "X2", pch = "+")  
3 points(x1[y == 1], x2[y == 1], col = "blue", pch = 4)
```



```
In [22]: 1 #plot show non-linear decision boundry
```

#c

```
In [23]: 1 lm.fit = glm(y~x1 + x2, family = binomial)
          2 summary(lm.fit)
```

Call:

```
glm(formula = y ~ x1 + x2, family = binomial)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.341	-1.212	1.044	1.124	1.253

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.10321	0.08976	1.150	0.250
x1	-0.40262	0.30981	-1.300	0.194
x2	-0.21758	0.30929	-0.703	0.482

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 691.79 on 499 degrees of freedom


Residual deviance: 689.57 on 497 degrees of freedom

AIC: 695.57

Number of Fisher Scoring iterations: 3

#d

```
In [24]: 1 data = data.frame(x1 = x1, x2 = x2, y = y)
2 lm.prob = predict(lm.fit, data, type = "response")
3 lm.pred = ifelse(lm.prob > 0.52, 1, 0)
4 data.pos = data[lm.pred == 1, ]
5 data.neg = data[lm.pred == 0, ]
6 plot(data.pos$x1, data.pos$x2, col = "blue", xlab = "X1",
7       ylab = "X2", pch = "+")
8 points(data.neg$x1, data.neg$x2, col = "red", pch = 4)
```



```
In [25]: 1 #no decision boundry is shown in the plot. all points are clas
```

#e

```
In [26]: 1 lm.fit = glm(y~poly(x1, 2) + poly(x2, 2) + I(x1 * x2),  
2          data = data, family = binomial)
```


Warning message:

"glm.fit: algorithm did not converge"Warning message:

"glm.fit: fitted probabilities numerically 0 or 1 occurred"

#f


```
In [27]: 1 lm.prob = predict(lm.fit, data, type = "response")
2         lm.pred = ifelse(lm.prob > 0.5, 1, 0)
3         data.pos = data[lm.pred == 1, ]
4         data.neg = data[lm.pred == 0, ]
5         plot(data.pos$x1, data.pos$x2, col = "blue", xlab = "X1",
6             ylab = "X2", pch = "+")
7         points(data.neg$x1, data.neg$x2, col = "red", pch = "x")
```



```
In [28]: 1 #non-linear decision boundary very similar to true decision bounda
```

#g


```
In [29]: 1 svm.fit = svm(as.factor(y) ~ x1 + x2, data, kernel = "linear",  
2               cost = 0.1)  
3 svm.pred = predict(svm.fit, data)  
4 data.pos = data[svm.pred == 1, ]  
5 data.neg = data[svm.pred == 0, ]  
6 plot(data.pos$x1, data.pos$x2, col = "blue", xlab = "X1",  
7       ylab = "X2",  
8       pch = "+")  
9 points(data.neg$x1, data.neg$x2, col = "red", pch = 4)
```



```
In [30]: 1 # linear kernel, even with low cost fails to find  
2 #non-linear decision boundary  
3 #and classifies all points to a single class.
```


#h

```
In [31]: 1 svm.fit = svm(as.factor(y)~x1 + x2, data, gamma = 1)
          2 svm.pred = predict(svm.fit, data)
          3 data.pos = data[svm.pred == 1, ]
          4 data.neg = data[svm.pred == 0, ]
          5 plot(data.pos$x1, data.pos$x2, col = "blue", xlab = "X1",
          6       ylab = "X2", pch = "+")
          7 points(data.neg$x1, data.neg$x2, col = "red", pch = 4)
```



```
In [32]: 1 #non-linear decision boundary on predicted labels similar
          2 #true decision boundary.
```

```
In [33]: 1 #from above problem, SVMs of non-linear kernel are extremely power
2 #in finding non-linear boundary.logistic regression with non-inter
3 #and SVMs with linear kernels fail to find the decision boundary.
4 #Adding interaction terms to logistic regression gives same power
5 #radial-basis kernels. However, there is some manual efforts and
6 #tuning involved in picking right interaction terms. This effort c
7 #prohibitive with large number of features. Radial basis kernels,
8 #on the other hand, only require tuning of one parameter - gamma -
9 #which can be easily done using cross-validation.
```

#7

```
In [34]: 1 library(ISLR)
2 gas.med = median(Auto$mpg)
3 new.var = ifelse(Auto$mpg > gas.med, 1, 0)
4 Auto$mpglevel = as.factor(new.var)
```

```
In [35]: 1 library(e1071)
```

```
In [36]: 1 set.seed(3255)
2 tune.out = tune(svm, mpglevel ~ ., data = Auto, kernel = "linear",
3               ranges = list(cost = c(0.01,
4               0.1, 1, 5, 10, 100)))
5 summary(tune.out)
```

Parameter tuning of 'svm':

– sampling method: 10-fold cross validation

– best parameters:

```
cost
1
```

– best performance: 0.01269231

– Detailed performance results:

```
cost      error dispersion
1 1e-02 0.07397436 0.06863413
2 1e-01 0.05102564 0.06923024
3 1e+00 0.01269231 0.02154160
4 5e+00 0.01519231 0.01760469
5 1e+01 0.02025641 0.02303772
6 1e+02 0.03294872 0.02898463
```

```
In [37]: 1 set.seed(21)
          2 tune.out = tune(svm, mpglevel ~ ., data = Auto, kernel = "polynomi
          3                      ranges = list(cost = c(0.1,
          4                      1, 5, 10), degree = c(2, 3, 4)))
          5 summary(tune.out)
```

Parameter tuning of 'svm':

– sampling method: 10-fold cross validation

– best parameters:

cost	degree
10	2

– best performance: 0.5435897

– Detailed performance results:

	cost	degree	error	dispersion
1	0.1	2	0.5587821	0.04538579
2	1.0	2	0.5587821	0.04538579
3	5.0	2	0.5587821	0.04538579
4	10.0	2	0.5435897	0.05611162
5	0.1	3	0.5587821	0.04538579
6	1.0	3	0.5587821	0.04538579
7	5.0	3	0.5587821	0.04538579
8	10.0	3	0.5587821	0.04538579
9	0.1	4	0.5587821	0.04538579
10	1.0	4	0.5587821	0.04538579
11	5.0	4	0.5587821	0.04538579
12	10.0	4	0.5587821	0.04538579

```
In [38]: 1 set.seed(333)
2 tune.out = tune(svm, mpglevel ~ ., data = Auto, kernel = "radial",
3               ranges = list(cost = c(0.1,
4               1, 5, 10), gamma = c(0.01, 0.1, 1, 5, 10, 100)))
5 summary(tune.out)
```

Parameter tuning of 'svm':

– sampling method: 10-fold cross validation

– best parameters:

```
cost gamma
10 0.01
```

– best performance: 0.02301282

– Detailed performance results:

	cost	gamma	error	dispersion
1	0.1	1e-02	0.08647436	0.05202669
2	1.0	1e-02	0.07397436	0.03896185
3	5.0	1e-02	0.04346154	0.03829236
4	10.0	1e-02	0.02301282	0.03298698
5	0.1	1e-01	0.07647436	0.04153447
6	1.0	1e-01	0.04858974	0.03714976
7	5.0	1e-01	0.02820513	0.04090081
8	10.0	1e-01	0.02564103	0.02702801
9	0.1	1e+00	0.55628205	0.04963230
10	1.0	1e+00	0.06628205	0.03615496
11	5.0	1e+00	0.05608974	0.03755733
12	10.0	1e+00	0.05608974	0.03755733
13	0.1	5e+00	0.55628205	0.04963230
14	1.0	5e+00	0.51025641	0.05558924
15	5.0	5e+00	0.50512821	0.06013429
16	10.0	5e+00	0.50512821	0.06013429
17	0.1	1e+01	0.55628205	0.04963230
18	1.0	1e+01	0.53076923	0.06373112
19	5.0	1e+01	0.52564103	0.06043648
20	10.0	1e+01	0.52564103	0.06043648
21	0.1	1e+02	0.55628205	0.04963230
22	1.0	1e+02	0.55628205	0.04963230
23	5.0	1e+02	0.55628205	0.04963230
24	10.0	1e+02	0.55628205	0.04963230


In [39]:

```


1 svm.linear = svm(mpglevel~., data = Auto, kernel = "linear",
2                 cost = 1)
3
4 svm.poly = svm(mpglevel~., data = Auto, kernel = "polynomial",
5               cost = 10,
6               degree = 2)
7 svm.radial = svm(mpglevel~., data = Auto, kernel = "radial",
8                 cost = 10, gamma = 0.01)
9 plotpairs = function(fit) {
10   for (name in names(Auto)[!(names(Auto)
11                             %in% c("mpg", "mpglevel", "name"))])
12     plot(fit, Auto, as.formula(paste("mpg~", name, sep = "")))
13 }
14 }
15 plotpairs(svm.linear)

```

SVM classification plot



```
In [40]: 1 plotpairs(svm.poly)
```



```
In [41]: 1 plotpairs(svm.radial)
```



```
In [ ]: 1
```

