


#Lab: Support Vector Machines

```
In [1]: 1 library(e1071)
```

```
In [2]: 1 set.seed(1)
2 x=matrix(rnorm(20*2),ncol=2)
3 y=c(rep(-1,10),rep(1,10))
4 x[y==1,]=x[y==1,]+1
5 plot(x,col=(3-y))
```



```


In [3]: 1 #svm() function can be used to fit a support vector classifier
        2 #when the argument kernel="linear" is used.
        3 dat=data.frame(x=x,y=as.factor(y)) #resp as factor variable
        4 library(e1071)
        5 svmfit=svm(y~., data=dat, kernel="linear", cost=10,
        6 scale=FALSE) #svm the support vector classifier for a given value
        7
        8 #False--not to scale each feature to have mean zero or standard de

```

```

In [4]: 1 plot(svmfit , dat) #plot support vector classifier

```



```

In [5]: 1 #-1 class is shown in yellow, and +1 class is shown in red,
        2 #disc boundry is "linear"

```

```
In [6]: 1 svmfit$index
        2
        3 #The support vectors are plotted as crosses and the remaining
        4 #observations are plotted as circles
        5 #even support vectors are below-

1  2  5  7 14 16 17
```

```
In [7]: 1 summary(svmfit)
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 10, scale
= FALSE)
```

Parameters:

```
SVM-Type:  C-classification
SVM-Kernel: linear
      cost:  10
      gamma: 0.5
```

Number of Support Vectors: 7

```
( 4 3 )
```


Number of Classes: 2

Levels:

```
-1 1
```

```
In [8]: 1 #linear kernel was used with cost=10, and that there were seven
        2 #support vectors, four in one class and three in the other.
```

```
In [9]: 1 #if we use small value of cost
2 svmfit=svm(y~., data=dat, kernel="linear", cost=0.1, scale=FALSE)
3 plot(svmfit , dat)
```



```
In [10]: 1 svmfit$inde
2
3 #smaller cost, larger number of support vectors, coz the margin is
```

1 2 3 4 5 7 9 10 12 13 14 15 16 17 18 20

```
In [11]: 1 set.seed (1)
          2 tune.out=tune(svm,y~.,data=dat,kernel="linear",
          3 ranges=list(cost=c(0.001, 0.01, 0.1, 1,5,10,100)))
          4 ##tune() performs ten-fold cross-validation on a set
          5 summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost
0.1

- best performance: 0.05

- Detailed performance results:

	cost	error	dispersion
1	1e-03	0.55	0.4377975
2	1e-02	0.55	0.4377975
3	1e-01	0.05	0.1581139
4	1e+00	0.15	0.2415229
5	5e+00	0.15	0.2415229
6	1e+01	0.15	0.2415229
7	1e+02	0.15	0.2415229

```
In [12]: 1 #cost=0.1 results in the lowest cross-validation error rate.
          2
```

```
In [13]: 1 bestmod=tune.out$best.model #tune() stores the best model obtained,
          2 summary(bestmod)
```

Call:

```
best.tune(method = svm, train.x = y ~ ., data = dat, ranges = list(cost = c(0.001,
                                0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 0.1
gamma: 0.5
```

Number of Support Vectors: 16

```
( 8 8 )
```

Number of Classes: 2

Levels:

```
-1 1
```

```
In [14]: 1 xtest=matrix(rnorm(20*2), ncol=2)
          2 ytest=sample(c(-1,1), 20, rep=TRUE)
          3 xtest[ytest==1,]=xtest[ytest==1,] + 1
          4 testdat=data.frame(x=xtest, y=as.factor(ytest))
          5
          6 #predict() can be used to predict the class label on a
          7 #set of test observations, at any given value of the cost parameter
          8
```

```
In [15]: 1 #predicting the class labels of these test observations.
          2 ypred=predict(bestmod,testdat)
          3 table(predict=ypred, truth=testdat$y)
```


```
      truth
predict -1 1
      -1  9 1
       1  2 8
```

```
In [16]: 1 svmfit=svm(y~.,data=dat, kernel="linear", cost=.01, scale=FALSE)
          2 ypred=predict(svmfit,testdat)
          3 table(predict=ypred,truth=testdat$y)
```


```
      truth
predict -1  1
      -1 11  6
       1  0  3
```

```
In [17]: 1 # three additional observation is misclassified.
```

```
In [18]: 1 #In case two classes are linearly separable.  
2 x[y==1,]=x[y==1,]+0.5  
3 plot(x, col=(y+5)/2, pch=19)
```




```
In [19]: 1 x[y==1,]=x[y==1,]+0.5  
2 plot(x, col=(y+5)/2, pch=19)
```



```
In [20]: 1 #observations are just barely linearly separable. We fit the SV ca  
2 #and plot the resulting hyperplane, using a very large value of co  
3 #so that no observations are misclassified
```

```
In [21]: 1 dat=data.frame(x=x,y=as.factor(y))
          2 svmfit=svm(y~., data=dat, kernel="linear",cost=1e5)
          3 summary(svmfit)
          4
          5 #No training errors were made and only three support vectors were
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 1e+05)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 1e+05
gamma: 0.5
```

Number of Support Vectors: 3

```
( 1 2 )
```

Number of Classes: 2

Levels:

```
-1 1
```

```
In [22]: 1 svmfit=svm(y~., data=dat, kernel="linear", cost=1)
          2 summary(svmfit)
          3 plot(svmfit ,dat)
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 1)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 1
gamma: 0.5
```

Number of Support Vectors: 5


```
( 2 3 )
```

Number of Classes: 2

Levels:

-1 1

SVM classification plot



In [56]:

```

1 #from fig, margin is very narrow as the observations that are not
2 #indicated as circles, are very close to the decision boundary.
3 #It seems likely that this model will perform poorly on test data.
4 #close to the decision boundary. It seems likely that this model

```

In [24]:

```

1 # trying with smaller value of cost
2 svmfit=svm(y~., data=dat, kernel="linear", cost=1)
3 summary(svmfit)

```

```
summary(svmfit)  
plot(svmfit, dat)
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 1)
```

Parameters:

```
SVM-Type: C-classification  
SVM-Kernel: linear  
cost: 1  
gamma: 0.5
```

Number of Support Vectors: 5


(2 3)


Number of Classes: 2

Levels:

-1 1

SVM classification plot





In [25]:

```

1 #with cost=1, we misclassify a training observation, but we obtain
2 #much wider margin and make use of seven support vectors.
3 # this model will perform better on test data than the model with

```

#Support Vector Machine


In [26]:

```

1 set.seed(3)
2 x=matrix(rnorm(200*2), ncol=2)
3 x[1:100,]=x[1:100,]+2
4 x[101:150,]=x[101:150,]-2
5 y=c(rep(1,150),rep(2,50))
6 dat=data.frame(x=x,y=as.factor(y))

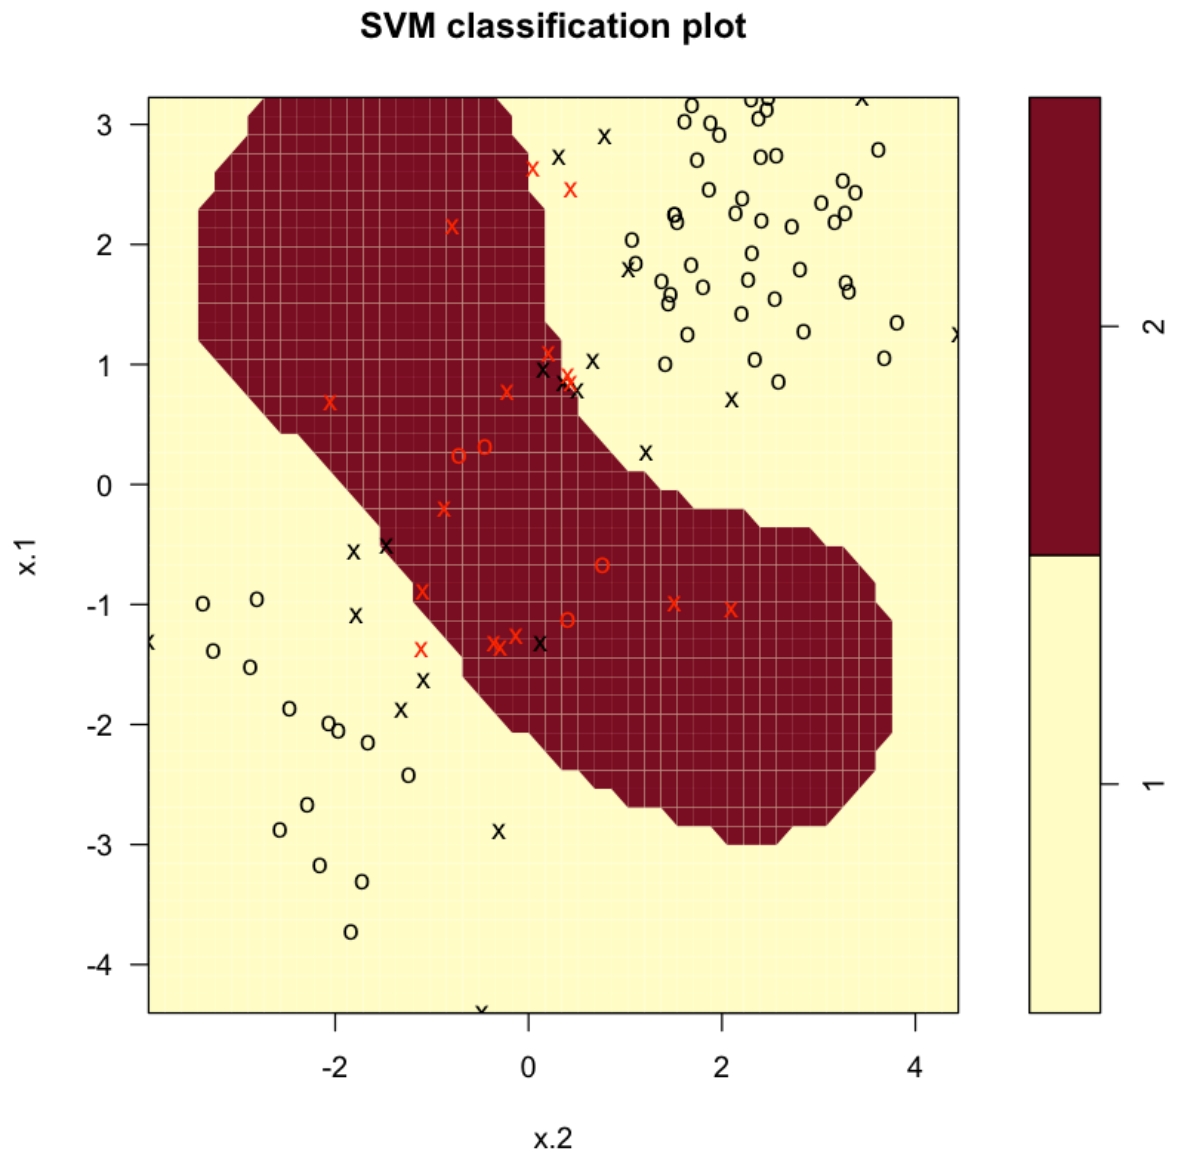
```

```
In [27]: 1 plot(x, col=y)
```



```
In [28]: 1 #data randomly split into training and testing groups.  
2 #fitting the training data using the svm()with a radial kernel and
```

```
In [29]: 1 train=sample(200,100)
          2 svmfit=svm(y~., data=dat[train,], kernel="radial", gamma=1,
          3 cost =1)
          4 plot(svmfit,dat[train,])
```



```
In [30]: 1 #rplot shows esulting SVM has a decidedly non-linear boundary.  
2 summary(svmfit)
```

Call:

```
svm(formula = y ~ ., data = dat[train, ], kernel = "radial", gamma =  
1,  
    cost = 1)
```

Parameters:

```
SVM-Type: C-classification  
SVM-Kernel: radial  
    cost: 1  
    gamma: 1
```

Number of Support Vectors: 36

```
( 20 16 )
```


Number of Classes: 2

Levels:

```
1 2
```



```
In [31]: 1 svmfit=svm(y~., data=dat[train,], kernel="radial",gamma=1, cost=1e
          2 plot(svmfit ,dat[train ,])
```



```
In [32]: 1 #fair number of training errors in this SVM fit from plot.
          2 #If we increase the value of cost, we can reduce the n0. of training
          3 #but there will be more irregular decision boundary that may overfit
```

```
In [33]: 1 set.seed (1)
          2 tune.out=tune(svm, y~., data=dat[train,], kernel="radial",
          3 ranges=list(cost=c(0.1,1,10,100,1000),
          4 gamma=c(0.5,1,2,3,4) ))
          5 summary(tune.out)
```

Parameter tuning of 'svm':

– sampling method: 10-fold cross validation

– best parameters:

cost	gamma
1	1

– best performance: 0.09

– Detailed performance results:

	cost	gamma	error	dispersion
1	1e-01	0.5	0.20	0.11547005
2	1e+00	0.5	0.11	0.08755950
3	1e+01	0.5	0.14	0.08432740
4	1e+02	0.5	0.11	0.05676462
5	1e+03	0.5	0.11	0.08755950
6	1e-01	1.0	0.20	0.11547005
7	1e+00	1.0	0.09	0.07378648
8	1e+01	1.0	0.11	0.07378648
9	1e+02	1.0	0.11	0.08755950
10	1e+03	1.0	0.12	0.10327956
11	1e-01	2.0	0.20	0.11547005
12	1e+00	2.0	0.11	0.07378648
13	1e+01	2.0	0.12	0.09189366
14	1e+02	2.0	0.15	0.09718253
15	1e+03	2.0	0.14	0.10749677
16	1e-01	3.0	0.20	0.11547005
17	1e+00	3.0	0.11	0.07378648
18	1e+01	3.0	0.14	0.08432740
19	1e+02	3.0	0.15	0.09718253
20	1e+03	3.0	0.13	0.08232726
21	1e-01	4.0	0.20	0.11547005
22	1e+00	4.0	0.12	0.07888106
23	1e+01	4.0	0.15	0.09718253
24	1e+02	4.0	0.17	0.10593499
25	1e+03	4.0	0.16	0.10749677

```
In [34]: 1 #the best choice of parameters involves cost=1 and gamma=1
```

In [35]:

```

1 #predict() to view test set predictions for this model
2
3 table(true=dat[-train,"y"], pred=predict(tune.out$best.model,
4                                           newdata=dat[-train ,]))

```

```

      pred
true 1  2
1 65  5
2  8 22

```

In [36]:

```

1 ### # ROC Curves
2 library(ROCR)
3

```

Loading required package: gplots

Attaching package: 'gplots'

The following object is masked from 'package:stats':

```
lowess
```

In [37]:

```

1 rocplot=function(pred, truth, ...){
2   predob = prediction (pred, truth)
3   perf = performance (predob , "tpr", "fpr")
4   plot(perf ,...)
5 }

```


In [38]:

```

1 svmfit.opt=svm(y~., data=dat[train,], kernel="radial", gamma=2,
2               cost=1,decision.values=T)
3
4 fitted=attributes(predict(svmfit.opt,dat[train,],
5                           decision.values=TRUE))$decision.values
6

```


```
In [39]: 1 #Roc plot
2 par(mfrow=c(1,2))
3 rocplot(fitted ,dat[train ,"y"],main="Training Data")
4
```



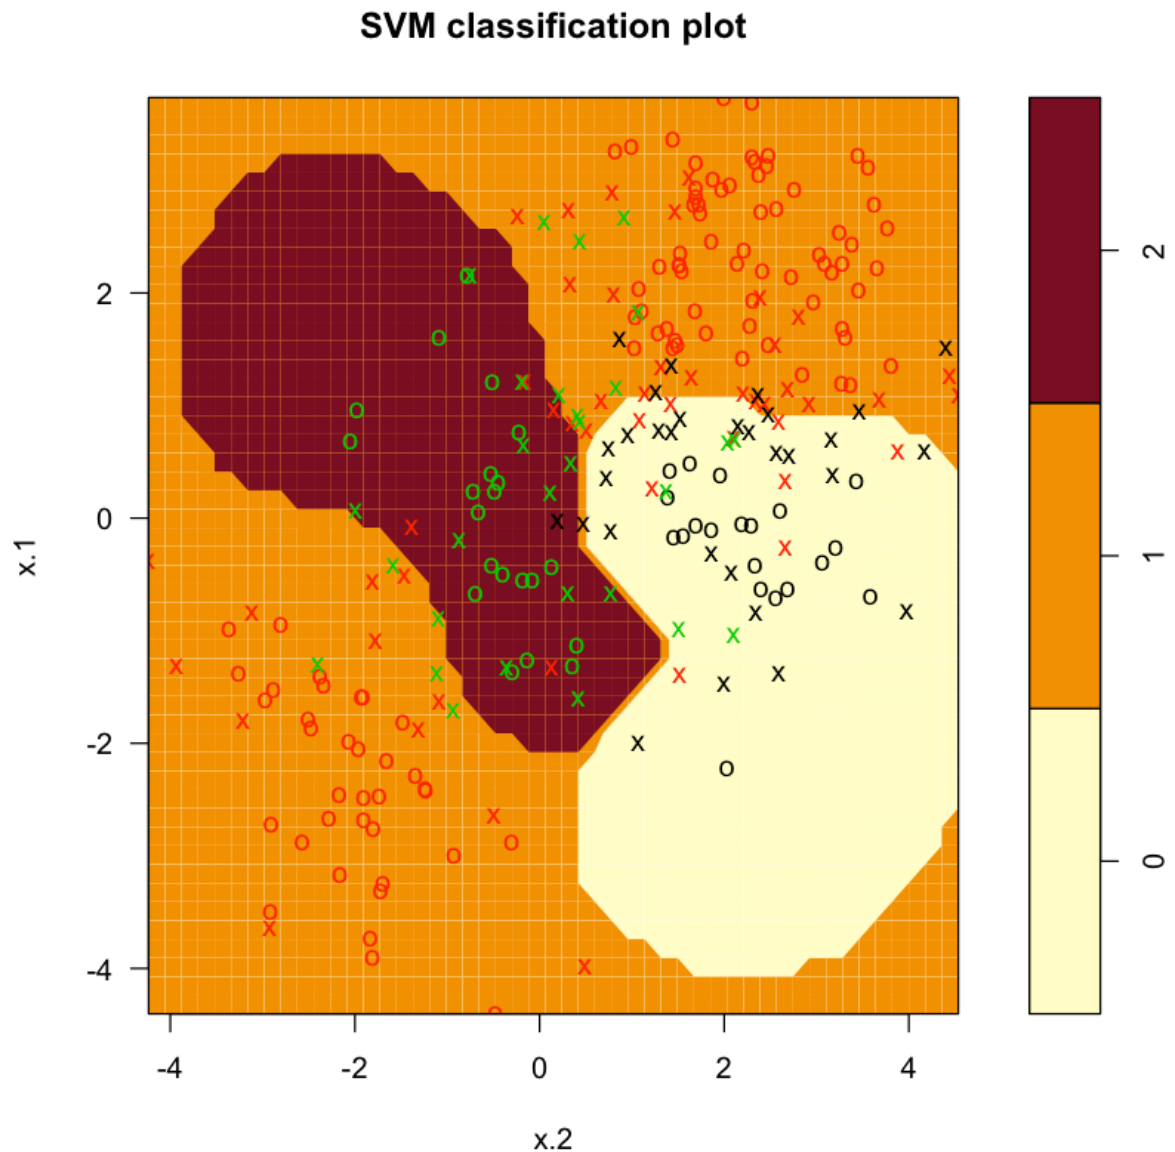
```
In [40]: 1 #SVM appears to be producing accurate predictions. increasing γ ca
2 #more flexible fit and generate further improvements in accuracy.
3
```

```
In [41]: 1 ### SVM with multiclass
2 #If the response is a factor containing more than two levels, then
3 #function will perform multi-class classification
```

```
In [42]: 1 set.seed(1)
2 x=rbind(x, matrix(rnorm(50*2), ncol=2))
3 y=c(y, rep(0,50))
4 x[y==0,2]=x[y==0,2]+2
5 dat=data.frame(x=x, y=as.factor(y))
6 par(mfrow=c(1,1))
7 plot(x,col=(y+1))
```



```
In [43]: 1 #fitting svm to data
          2 svmfit=svm(y~., data=dat, kernel="radial", cost=10, gamma=1)
          3 plot(svmfit , dat)
```



```
In [44]: 1 #the e1071 library can also be used to perform sv regression,
          2 #if the resp vector
          3 #passed in to svm() is numerical rather than a factor.
```

```
In [45]: 1 #Application to Gene Expression Data
          2
          3 library(ISLR)
          4 names(Khan)
```

```
'xtrain' 'xtest' 'ytrain' 'ytest'
```

```
In [46]: 1 dim(Khan$xtrain)
```

```
63 2308
```

```
In [47]: 1 dim(Khan$xtest )
```

```
20 2308
```

```
In [48]: 1 length(Khan$ytrain )
```

```
63
```

```
In [49]: 1 length(Khan$ytest)
```

```
20
```

```
In [50]: 1 table(Khan$ytrain)
```

```
1 2 3 4
8 23 12 20
```

```
In [51]: 1 table(Khan$ytest)
```

```
1 2 3 4
3 6 6 5
```

```
In [52]: 1 dat=data.frame(x=Khan$xtrain , y=as.factor(Khan$ytrain ))
          2 out=svm(y~., data=dat, kernel="linear",cost=10)
          3 summary(out)
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 10)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 10
gamma: 0.0004332756
```

Number of Support Vectors: 58

```
( 20 20 11 7 )
```

Number of Classes: 4

Levels:

```
1 2 3 4
```

```
In [53]: 1 table(out$fitted , dat$y)
```

```
      1  2  3  4
1  8  0  0  0
2  0 23  0  0
3  0  0 12  0
4  0  0  0 20
```

```
In [54]: 1 dat.te=data.frame(x=Khan$xtest , y=as.factor(Khan$ytest ))
          2 pred.te=predict(out, newdata=dat.te)
          3 table(pred.te, dat.te$y)
```

```
pred.te 1 2 3 4
      1 3 0 0 0
      2 0 6 2 0
      3 0 0 4 0
      4 0 0 0 5
```



```
In [55]: 1 #there are no training errors.because the large number of variable  
2 #relative to the no. of observations implies that it is easy to fi  
3 #hyperplanes that fully separate the classes. We are most interest  
4 #in the SV classifier's performance on the training observations,  
5 #but rather its performance on the test observations
```

```
In [ ]: 1
```