

M6 applied**5(a)**

```
In [28]: 1 library(ISLR)
```

```
In [29]: 1 data(Default)
```

```
In [30]: 1 set.seed(1)
```

```
In [58]: 1 fit1 = glm(default~income + balance, data=Default,
          2          family=binomial)
```

```
In [32]: 1 summary(fit1)
```

Call:

```
glm(formula = default ~ income + balance, family = binomial,
    data = Default)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4725	-0.1444	-0.0574	-0.0211	3.7245

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.154e+01	4.348e-01	-26.545	< 2e-16 ***
income	2.081e-05	4.985e-06	4.174	2.99e-05 ***
balance	5.647e-03	2.274e-04	24.836	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom
 Residual deviance: 1579.0 on 9997 degrees of freedom
 AIC: 1585

Number of Fisher Scoring iterations: 8

#5(b)

```
In [33]: 1 set.seed(1)
          2 train = sample(nrow(Default), nrow(Default)*0.5)
```

```
In [34]: 1 fit2 = glm(default~income+balance, data=Default, family=binomial,
```

```
In [35]: 1 prob2 = predict(fit2, Default[-train,], type="response")
          2 pred2 = ifelse(prob2 > 0.5, "Yes", "No")
          3 table(pred2, Default[-train,]$default)
```

```
pred2  No  Yes
No    4824 108
Yes    19  49
```

```
In [36]: 1 mean(Default[-train,]$default != pred2)
```

```
0.0254
```

```
In [37]: 1 #test error is 0.0254
```

5(c)

```
In [59]: 1 #repeat 1
          2 set.seed(2)
          3 train = sample(nrow(Default), nrow(Default)*0.5)
          4 fit2 = glm(default~income+balance, data=Default, family=binomial,
          5               subset=train)
          6 prob2 = predict(fit2, Default[-train,], type="response")
          7 pred2 = ifelse(prob2 > 0.5, "Yes", "No")
          8 mean(Default[-train,]$default != pred2)
```

```
0.0238
```

```
In [60]: 1 #repeat 2
          2 train=sample(nrow(Default), nrow(Default)*0.5)
          3 fit2=glm(default~income+balance, data=Default, family=binomial,
          4               subset=train)
          5 prob2=predict(fit2, Default[-train,], type="response")
          6 pred2=ifelse(prob2 > 0.5, "Yes", "No")
          7 mean(Default[-train,]$default != pred2)
```

```
0.0288
```

```
In [61]: 1 train=sample(nrow(Default), nrow(Default)*0.5)
          2 fit2=glm(default~income+balance, data=Default, family=binomial,
          3             subset=train)
          4 prob2=predict(fit2, Default[-train,], type="response")
          5 pred2=ifelse(prob2 > 0.5, "Yes", "No")
          6 mean(Default[-train,]$default != pred2)
```

0.0254

```
In [41]: 1 # from three repetitions, the test error is 2.5% consistance
          2 #(0.02 appx)very less variance
```

#5(d)

```
In [62]: 1 set.seed(1)
          2 train=sample(nrow(Default), nrow(Default)*0.5)
          3 fit3=glm(default~income + balance + student, data=Default,
          4             family=binomial, subset=train)
          5 prob3=predict(fit3, Default[-train,], type="response")
          6 pred3=ifelse(prob3 > 0.5, "Yes", "No")
          7 mean(Default[-train,]$default != pred3)
```

0.026

```
In [43]: 1 #the test error is almost similar with adding "student".
          2 #there is no significance reduction in the error.
```

#6(a)

```
In [63]: 1 require(ISLR)
          2 data(Default)
          3 set.seed(1)
          4 fit1=glm(default~income+balance, data=Default,
          5           family=binomial)
          6 summary(fit1)
```

Call:

```
glm(formula = default ~ income + balance, family = binomial,
    data = Default)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4725	-0.1444	-0.0574	-0.0211	3.7245

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.154e+01	4.348e-01	-26.545	< 2e-16 ***
income	2.081e-05	4.985e-06	4.174	2.99e-05 ***
balance	5.647e-03	2.274e-04	24.836	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom
 Residual deviance: 1579.0 on 9997 degrees of freedom
 AIC: 1585

Number of Fisher Scoring iterations: 8

```
In [45]: 1 #estimated std error for income is 0.000004985 and for
          2 #balance is 0.0002274
```

6(b)

```
In [46]: 1 set.seed(1)
          2 boot.fn=function(df, trainid) {
          3   return(coef(glm(default~income + balance, data=df,
          4                       family=binomial, subset=trainid)))
          5 }
          6 boot.fn(Default, 1:nrow(Default))
```

```
(Intercept) -11.5404684366776
income      2.08089755005988e-05
balance     0.00564710294333903
```

6(c)

```
In [47]: 1 require(boot)
          2 boot(Default,boot.fn,R=100)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = Default, statistic = boot.fn, R = 100)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	-1.154047e+01	8.556378e-03	4.122015e-01
t2*	2.080898e-05	-3.993598e-07	4.186088e-06
t3*	5.647103e-03	-4.116657e-06	2.226242e-04


```
In [48]: 1 #std error estimates are very close using glm and bootstrap
          2 #using R=100
          3 #glm-(4.985e-06 for income,2.274e-04 for balance)
          4 #bootstrap-(4.186088e-06 for income,2.226242e-04 for balance)
```

8(a)

```
In [49]: 1 #n=100 (observations)
          2 #p=2 (features)
          3 #equation is--
          4 #  $Y = X - 2X^2 + \epsilon$ 
          5
          6 set.seed(1)
          7 x=rnorm(100)
          8 y=x-2*x^2 + rnorm(100)
```

8(b)

```
In [50]: 1 plot(x,y)
```



```
In [51]: 1 #relation b/w x and y is quadratic
```

#8(c)

```
In [52]: 1 set.seed(1)
2 df = data.frame(y, x, x2=x^2, x3=x^3, x4=x^4) #df dataframe
3 fit1 = glm(y~x, data=df)
4 cv.err1 = cv.glm(df, fit1)
5 cv.err1$delta
6 fit2 = glm(y~x + x2, data=df)
7 cv.err2 = cv.glm(df, fit2)
8 cv.err2$delta
9 fit3 = glm(y~x + x2 + x3, data=df)
10 cv.err3 = cv.glm(df, fit3)
11 cv.err3$delta
12 fit4 = glm(y~x + x2 + x3 + x4, data=df)
13 cv.err4 = cv.glm(df, fit4)
14 cv.err4$delta
```

7.28816160667281 7.28474411546929

0.937423637615552 0.937178917181123

0.956621830108939 0.956253813731321

0.953904892744804 0.953445283156601

#8(d)


```
In [53]: 1 set.seed(2)
2 df = data.frame(y, x, x2=x^2, x3=x^3, x4=x^4)
3 fit1 = glm(y~x, data=df)
4 cv.err1 = cv.glm(df, fit1)
5 cv.err1$delta
6 fit2 = glm(y~x + x2, data=df)
7 cv.err2 = cv.glm(df, fit2)
8 cv.err2$delta
9 fit3 = glm(y~x + x2 + x3, data=df)
10 cv.err3 = cv.glm(df, fit3)
11 cv.err3$delta
12 fit4 = glm(y~x + x2 + x3 + x4, data=df)
13 cv.err4 = cv.glm(df, fit4)
14 cv.err4$delta
```

7.28816160667281 7.2847441154693

0.937423637615552 0.937178917181122

0.95662183010894 0.956253813731321

0.953904892744804 0.953445283156601

```
In [54]: 1 # no diff in results b/w (c) and (d)
```

8(e)

```
In [55]: 1 #using quad model,x x^2 x^3 having the least error.
2 #it was expected as true model was generated by quadratic model
```

```
In [56]: 1 fitn = lm(y~poly(x,4))
          2 summary(fitn)
```

Call:

```
lm(formula = y ~ poly(x, 4))
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.0550	-0.6212	-0.1567	0.5952	2.2267

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.55002	0.09591	-16.162	< 2e-16 ***
poly(x, 4)1	6.18883	0.95905	6.453	4.59e-09 ***
poly(x, 4)2	-23.94830	0.95905	-24.971	< 2e-16 ***
poly(x, 4)3	0.26411	0.95905	0.275	0.784
poly(x, 4)4	1.25710	0.95905	1.311	0.193

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9591 on 95 degrees of freedom

Multiple R-squared: 0.8753, Adjusted R-squared: 0.8701

F-statistic: 166.7 on 4 and 95 DF, p-value: < 2.2e-16

```
In [57]: 1 #from summary, x and x^2 are the significant predictors.
          2 #this agrees with cross-validation results which indicates using
          3 #x and x^2 gives the best outcome.
```